

Self-Evaluation Exercise on OpenMP

1. Write three ways in which the total number of threads can be specified in an OpenMP program.

2. Consider these two codes. Are they equivalent?

Code A:

```
double res[MAX]; int i;
#pragma omp parallel
{
    #pragma omp for
    for (i=0; i< MAX; i++) {
        res[i] = huge();
    }
}
```

Code B:

```
double res[MAX]; int i;
#pragma omp parallel for
for (i=0; i< MAX; i++) {
    res[i] = huge();
}
```

3. Consider this piece of code.

```
const int MAX=6;
int big(int k) {return k;}
int main() {
```

```

int i, j, A[MAX];
j = 5;
for (i=0; i< MAX; i++) {
j +=2;
A[i] = big(j);
}
}

```

Can this be parallelized simply by writing "#pragma omp parallel for" just before the for loop? If yes, show the final value of array A. If no, change the code suitably to allow for-loop parallelization and then, show the final value of array A.

4. Replace WRITE_PRAGMA_HERE by suitable pragma to correctly compute average value in parallel.

```

double avg=0.0, A[MAX]; int i;
WRITE_PRAGMA_HERE
for (i=0; i< MAX; i++) {
avg += A[i];
}
avg = avg/MAX;

```

5. Consider “Serial_ComputingPi_IntegrationApproximation.cpp” program which shows serial code. The code computes pi by approximating integration operation.

Parallelize the code by adding suitable OpenMP directive.

Side Note: The pi value can also be computed Monte Carlo approach. See the following link: <http://jakascorner.com/blog/2016/05/omp-monte-carlo-pi.html>

6. In CNNs, the computation pattern of convolution layers and fully-connected (FC) layers can be modeled as matrix-matrix multiplication (MMM) and matrix-vector multiplication (MVM), respectively. Since MVM is just a simple case of MMM, we will now focus on MMM. As we will discuss later during this course, in AlexNet, convolution layers and FC layers account for 93\% and ~7% of total computations. Thus, any improvement in performance of convolution layer can directly improve the performance of AlexNet (and other CNNs in general).

See "Serial_MatMul.cpp". Parallelize it using OpenMP. Ensure that you clearly specify private/shared scope for all the variables. Observe speedup using different number of threads.

7. Consider the following code to compute pi. If required, replace WRITE_PRAGMA_HERE with suitable pragma to compute pi correctly. Otherwise, just say "no code required".

```
#include <omp.h>

#include <stdio.h>

static long num_steps = 10000000; double step;

#define NUM_THREADS 2

int main ()
{ double pi; step = 1.0/(double) num_steps;
  omp_set_num_threads(NUM_THREADS);
  int nthreads;
  #pragma omp parallel
  {
    int i, id,nthrds; double x, sum;
    id = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    if (id == 0) nthreads = nthrds;
    id = omp_get_thread_num();
    nthrds = omp_get_num_threads();
```

```
for (i=id, sum=0.0; i< num_steps; i=i+nthreads){  
    x = (i+0.5)*step;  
    sum += 4.0/(1.0+x*x);  
}  
//WRITE_PRAGMA_HERE  
pi += sum * step;  
}  
  
printf("pi = %f\n", pi);  
}
```